

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import tools
import warnings, sys
if not sys.warnoptions:
    warnings.simplefilter("ignore")
```

```
In [2]: # Load data
load_file = '../datasets/reddit_submissions.json'
someposts = pd.read_json(load_file, lines=True)
someposts.index = someposts['id']
```

## Model

Possible models:

- Bag of Words/Bigrams + LR/SVM
- Average Embedding + LR
- LDA
- Tree Kernels
- @RNN: I try word-based 1-layer LSTM as baseline
- @CNN: I secondly try character-level CNN
- @RCNN: I finally will try the advanced RCNN

I first try the simplest word-based LSTM, then try character-level CNN and finally try advanced word-based RCNN if possible.

```
In [8]: from tensorflow import keras
import tensorflow as tf

from keras.models import Sequential
from keras.layers import Dense, Embedding, LSTM

from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
from sklearn.model_selection import train_test_split
from keras.layers import Reshape, Conv1D, GlobalMaxPooling1D, MaxPooling1D

from sklearn.metrics import classification_report
```

## 1: word-based one-layer LSTM

LSTM could capture the contextual feature of text. Shown in Part1\_ModelTraining with avg f1-score is 0.62. This model is good to classify 3-'EatingDisorders' with 0.80 and 7-'alcoholism' with 0.81.

## 2: charactor-based one-layer CNN

The character-level cnn is effective and unbiased, but we have to define fixed size kernel while words with different length.

```
In [4]: maxlen = 2000
        vocabulary_size = 40

        # Preprocesses the raw data as required
        subreddit_mappings, someposts = tools.record_process(someposts, features

        # Fit the tokenizer on the training data
        t = Tokenizer(num_words=vocabulary_size, char_level=True)
        t.fit_on_texts(someposts[['title','selftext']].apply(lambda x: 'End. '.join(x)))
        print('%s unique words' %len(t.word_index))

        # partitions the model-ready data into train, validation, and test sets.
        X_train = t.texts_to_sequences(someposts[['title','selftext']].apply(lambda x: 'End. '.join(x)))
        X_train = pad_sequences(X_train, maxlen=maxlen, padding='post')
        X_train, X_val, y_train, y_val = train_test_split(X_train,someposts['target'])
        X_val, X_test, y_val, y_test = train_test_split(X_val,y_val, test_size=0.1)
        print('There are {}, {}, {} records for train, validation, and test sets'.format(len(X_train), len(X_val), len(X_test)))

        There are 236742 records after processing
        The sum of rare categories is 1142
        591 unique words
        There are 189393,23674,23675 records for train, validation, and test sets
```

```
In [12]: # training
n_classes = len(subreddit_mappings)

model = Sequential()
model.add(Reshape((maxlen, 1), input_shape=(maxlen,)))
model.add(Conv1D(150, 15, activation='relu'))
model.add(MaxPooling1D(pool_size=10, strides=5))
model.add(Conv1D(300, 5, activation='relu'))
model.add(GlobalMaxPooling1D())
model.add(Dense(n_classes, activation='softmax'))
model.summary()

# compile
model.compile('adam', 'sparse_categorical_crossentropy', metrics=['accuracy'])
```

Layer (type)	Output Shape	Param #
reshape_5 (Reshape)	(None, 2000, 1)	0
conv1d_6 (Conv1D)	(None, 1986, 150)	2400
max_pooling1d_5 (MaxPooling1D)	(None, 396, 150)	0
conv1d_7 (Conv1D)	(None, 392, 300)	225300
global_max_pooling1d_2 (GlobalMaxPooling1D)	(None, 300)	0
dense_3 (Dense)	(None, 17)	5117
Total params: 232,817		
Trainable params: 232,817		
Non-trainable params: 0		

```

In [*]: # prepare for training
early_stopping = keras.callbacks.EarlyStopping(monitor='acc',
                                                min_delta=0.0001,
                                                patience=1,
                                                verbose=1)

checkpoint = keras.callbacks.ModelCheckpoint('ModelTraining/lstm_1st.hdf5',
                                             verbose=1,
                                             save_best_only=True)

# training
history = model.fit(X_train, y_train,
                    batch_size = 64,
                    epochs=5,
                    validation_data=(X_val, y_val),
                    callbacks=[early_stopping])

# plot the training process
acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(len(acc))

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss(fine-tuning)')
plt.legend()

plt.show()

```

Train on 189393 samples, validate on 23674 samples

Epoch 1/5

70720/189393 [=====>.....] - ETA: 31:23 - loss: 2.0439 - acc: 0.3816

```

In [ ]: # Evaluation
results = model.predict(X_test)
predictions = results.argmax(axis = 1)
print(classification_report(y_test, predictions))

```

### 3: recurrent convolutional neural network

According to paper, RecurrentNN is a biased model, where later words are more dominant than earlier words. The Convolutional Neural Network (CNN) is unbiased to NLP tasks, which can fairly determine discriminative phrases in a text with a max-pooling layer, but it is difficult to determine the window size. Hence, RCNN for text classification that learn more contextual information than conventional window-based neural networks and represent the semantic of texts more precisely for text classification.

'First, we apply a bi-directional recurrent structure. Second, we employ a max-pooling layer that automatically judges which features play key roles in text classification, to capture the key component in the texts. '

For the '2015-fancy' model-RCNN. Under time limitation, I plan to use other codes as reference and fit it to our model.

- The paper: Recurrent Convolutional Neural Network for Text Classification (AAAI 2015)  
<https://www.aaai.org/ocs/index.php/AAAI/AAAI15/paper/view/9745>  
(<https://www.aaai.org/ocs/index.php/AAAI/AAAI15/paper/view/9745>)
- The reference code: <https://github.com/roomylee/rcnn-text-classification>  
(<https://github.com/roomylee/rcnn-text-classification>); <https://github.com/sklan/rcnn-text-classification> (<https://github.com/sklan/rcnn-text-classification>)

In [ ]: